

The Human in the Loop

Flyxion

September 2026

A person sits at the end of the pipeline. Their title contains the word "review." Their presence is cited, whenever the system is questioned, as the reason its output can be trusted: nothing ships, nothing decides, nothing acts, without a human confirming it first.

The confirmation is given an average of several seconds. The interface presents a recommendation, already formatted, already framed as the likely answer, alongside a button that accepts it and a much longer, much less convenient path that overrides it. The volume of decisions routed past this person is calibrated to the volume a person can process only by approving nearly all of them, because a review process designed to actually catch errors at scale would require a staff and a schedule the deployment was built specifically to avoid needing. The person is real. The loop is real. The judgment the arrangement is named for is not, because judgment requires time and information neither of which the design provides, and a design that withholds both from the one party it calls the decision-maker has not distributed responsibility to that party. It has stationed them there.

The function this arrangement serves becomes clear only once something goes wrong, which is the same moment every mechanism in this series reveals what it was built for. The automated system that made the actual recommendation is not named in the incident report as the responsible party, because a recommendation is not a decision. The human who clicked accept is named, because a click is a decision, and a decision has a decider. The company that built a review interface incapable of supporting review gets to describe itself as an organization with human oversight built in, right up until the oversight is asked to answer for an outcome it was never given the material conditions to actually prevent. The ceremony was the point. It was never meant to catch the error. It was meant to have someone standing next to it when it was found.