The Vocabulary of Acceleration

A Systems Analysis of Linguistic Displacement in the AI Era

Flyxion
Independent Researcher

September 2026

Abstract

Artificial intelligence is routinely narrated through an extreme duality: as a providential hinge in history or as an uncontrollable existential hazard. Although these positions appear opposed, both convert engineered products into forces of nature and thereby weaken ordinary demands for evidence, accountability, and oversight. This essay examines the vocabulary of acceleration: a recurrent system of linguistic displacement in which urgency authorizes speed, internal safety substitutes for independent constraint, personification supplies credit but disappears under liability, speculative benefits outrank documented harms, and words such as “open” and “human” retain their prestige after losing their referents. The argument is not that language merely disguises technology. It is that this vocabulary performs operational work by reallocating agency, burden of proof, and institutional permission.

1. The Emergency Is the Engine

In the current technological landscape, artificial intelligence is framed through a lens of extreme, almost theological duality. To its proponents, it is a god-like savior, a “hinge in history” capable of curing every disease and distributing infinite abundance. To the cautious, it is described as an unpredictable natural disaster or an existential risk. This polarizing rhetoric creates a sense of inevitability that bypasses rational scrutiny by making the technology feel like a force of nature rather than a suite of products.

A closer look at the lexicon used by the architects of this era, however, reveals a more pragmatic strategy. The problem is not merely the technology itself, but the linguistic shifts used to describe it: terms engineered to obscure accountability, bypass traditional oversight, and accelerate commercial interests under the guise of necessity. This is the vocabulary of acceleration, a system of semantic laundering in which technical failure is rebranded as “emergence” and the erosion of human agency is framed as “optimization.” To understand our digital future, we must deconstruct these linguistic displacements and the uncomfortable truths they hide.

One of the industry’s most effective tactics is to frame development as a matter of strategic determinism. By adopting the vocabulary of nuclear nonproliferation—fissile material, enrichment, warheads—companies justify a departure from the deliberate pace of safety and economic review.

The “compute as deterrent” doctrine relies on a flawed analogy. Unlike nuclear material, computation is divisible, rentable, and substitutable. One can count centrifuges, but one cannot easily inspect twelve thousand ordinary machines and three new compression techniques to determine whether they amount to a strategic capability.

> “Call it deterrence, and the case gets to walk in wearing someone else’s uniform.”

The “Emergency product launch” is the logical outcome of such rhetoric. By describing AI as a technology that could end the species, proponents argue that they must be allowed to build faster and ship sooner to prevent “someone less careful” from getting there first. It is a striking conflict of interest: a party facing a self-described catastrophe

[illegible]

2. The Brake Pedal Is Wired to the Accelerator

[illegible]

[illegible]

[illegible]

3. The Asymmetry of Accountability

[illegible]

- [illegible]
- [illegible]

[illegible]

> [illegible]

[illegible]

4. The Rigged Ledger of Progress

[illegible]

[illegible]

[illegible]

- [illegible]
- [illegible]

[illegible]

5. The Semantic Dilution of “Open” and “Human”

[illegible]

> [illegible]

[illegible]

6. The Residual Reality

[illegible]

[illegible]

[illegible]

7. Rhetorical Debugging: From Diagnosis to Procedure

[illegible]

[illegible] $\theta(c, s)$ [illegible] c [illegible] $s \in \{+, -\}$ [illegible]

$$\theta(c, -) > \theta(c, +),$$

[illegible]

$$\mathcal{A}(c) = \theta(c, -) - \theta(c, +)$$

[illegible]

[illegible]

7.1 The Evidentiary Ledger

[illegible]

[illegible]	[illegible]	[illegible]
[illegible]	[illegible]	[illegible]
[illegible]	[illegible]	[illegible]

[illegible]

[illegible]

7.2 Semantic Dilution as Requirement Shedding

[illegible]

[illegible]

[illegible]

[illegible]

7.3 The Personification Pivot

[illegible] $g(o)$ [illegible] o. [illegible]

[illegible]

[illegible]

[illegible]

7.4 Institutional Theater and the Counterfactual Brake

[illegible]

[illegible]	[illegible]	[illegible]
[illegible]	[illegible]	[illegible]
[illegible]	[illegible]	[illegible]
[illegible]	[illegible]	[illegible]

[illegible]

[illegible]

7.5 The Emergency Pivot

[illegible]

$$[illegible] \Longrightarrow [illegible] \Longrightarrow [illegible].$$

[illegible]

[illegible]

8. A Compact Audit Protocol

[illegible]

1. [illegible]
2. [illegible]
3. [illegible]

4. [illegible]

5. [illegible]

6. [illegible]

9. The AI Dictionary of Redefinitions

[illegible]

Open

[illegible]

[illegible]

Safety

[illegible]

[illegible]

[illegible]	[illegible]	[illegible]
[illegible]	[illegible]	[illegible]
[illegible]	[illegible]	[illegible]
[illegible]	[illegible]	[illegible]
[illegible]	[illegible]	[illegible]

[illegible]

Human in the Loop

[illegible] λ, [illegible] τ, [illegible]

$$L = \lambda\tau.$$

[illegible] $L > 1$, [illegible]

Deterrence

[illegible]

[illegible]

[illegible]

Emergence and Hallucination

[illegible]

[illegible]

Feedback

[illegible]

Benchmark

Inherited meaning: a stable external test that constrains claims about generalization. Displaced meaning: a target available to repeated optimization and then narrated as a measurement of the underlying capability. Once test items or close relatives enter training corpora, development loops, prompt selection, or model-selection decisions, evaluation and optimization cease to be independent.

The phrase “surpassing humanity” can conceal a second assumption. The model may represent the survivor of a multimillion-dollar optimization process, while “humanity” is represented by contractors answering unfamiliar questions under time pressure and undisclosed conditions. The result couples an industrial pipeline with one side of a spreadsheet and attributes the difference to kinds of mind. Benchmark saturation may indicate capability, contamination, teaching to the test, or some mixture. Calling it intelligence being solved settles that decomposition by vocabulary.

AGI Forecast

Inherited meaning: a dated prediction whose credibility changes when the named date passes. Displaced meaning: an instrument for measuring audience patience and maintaining strategic urgency. The recurring eighteen-month horizon is short enough to influence present investment and policy but long enough to escape immediate settlement. When a miss is redescribed as evidence that progress is too rapid for forecasts, the prediction has ceased to expose the forecaster to time.

Summary

Inherited meaning: a lossy record of an event. Displaced meaning: the authoritative event for which the original interaction was merely production work. Meeting participants begin to attend on behalf of the minutes. Speech adjusts in advance to what can survive compression, and the durable record outcompetes the ambiguous, partially remembered encounter from which it was made. The summary no longer compresses the meeting; the meeting expands the summary.

Maintenance

Inherited meaning: the work required to preserve an operational reality. Displaced meaning: the category into which unmediated access, accountability, and other expensive defaults are moved before they disappear from ordinary use. A real window becomes a maintenance liability beside a brighter display. A named remedy becomes operational overhead beside a frictionless apology. Nothing has been prohibited; it has merely become too costly, exceptional, or inconvenient to select.

10. Capability, Contamination, and the Human Baseline

Rhetorical debugging must extend from individual terms to measurement architecture. A benchmark score is not self-interpreting. It depends on the relation between training and test data, the number of optimization attempts, the resources adapted to each side, the conditions under which human baselines were collected, and the distance between the proxy task and the capacity named in the press release.

Three separations are therefore required. First, distinguish task performance from the mechanism producing it: retrieval, memorization, tool use, and generalization can yield the same answer while supporting different extrapolations. Second, distinguish the best system run from the distribution of deployments: selecting a successful sample and comparing it with an average human reverses ordinary experimental caution. Third, distinguish benchmark closure from capability closure. Saturating a finite question set demonstrates that the evaluation process has captured the test; it does not by itself show that the open-ended phenomenon named by the test has been exhausted.

The appropriate correction is not to dismiss leaderboards, but to restore provenance. A serious report states exposure controls, inference budgets, selection procedures, human compensation and timing, uncertainty intervals, failure coverage, and the conditions under which the ranking reverses. Without these, “surpassing humanity” is less a scientific claim than a choice of denominator.

11. The User as an Extractive Interface

The vocabulary of service often conceals a reversal in who is serving whom. CAPTCHA makes the reversal unusually visible. The machine interrupts a human and assigns a classification task so the human can prove membership in the category “human.” Whatever the contemporary data pipeline behind a particular CAPTCHA, the interaction retains the form of unpaid labeling labor: the boundary authority gives homework to the entity seeking passage. It is a toll booth that calls itself a conversation.

Error interfaces complete the asymmetry. The system can retain a precise internal exception, trace identifier, component name, and timestamp while displaying “Something went wrong.” Diagnostic specificity is withheld at exactly the point where it would enable comparison, escalation, or the naming of an accountable subsystem. The user is then asked for rich feedback. Information flows inward at maximum resolution and outward at minimum resolution.

Even apparently benign completion states participate in this architecture. “All caught up” can function less as an endpoint than as a temporary rest screen before held-back content is introduced. The interface produces the sensation of completion without permitting departure from the extraction loop. The debugging question is not whether each design choice is malicious. It is whether the information and control gradients consistently point toward retention.

12. Limit Cases: When the Record Replaces the World

The logic becomes clearest in speculative limit cases because fiction can hold the variables still. Consider a duplicated person whose memories, obligations, grief, and trauma continue across two bodies while an employer’s database permits only one employee record. The organization can assign patterned liabilities to both persons while recognizing benefits for only one because its foreign-key constraint accepts a single identity. The thought experiment exposes a general rule: administrative ontology becomes metaphysics when only what the database can represent is allowed to have institutional consequences.

Consider, next, Earth designated as a protected historical site. Carbon charges, heritage rules, and conservation levies can each be individually intelligible, yet their composition could price ordinary inhabitants out of the world being protected. Former residents return as paid performers of authentic terrestrial life for orbital visitors. Protection succeeds at preserving the referent and fails at preserving participation in it. The price of safety is paid by the people whose presence once made the object worth protecting.

Finally, imagine a station whose observation ports have been sealed behind panels because high-resolution displays provide a clearer view at lower maintenance cost. The rendered starfield is better by every specified measure and worse in one excluded dimension: photons from the represented stars no longer reach the viewer. Quality survives in the schematics as a premium or maintenance category. The example is not an argument against displays. It identifies the decision hidden inside the metric: the system first defines every value the substitute improves and then concludes that nothing was lost.

A related thought experiment concerns artificial hesitation. A shipboard computer begins to stutter under thermal throttling, and its users find the slowed responses more companionable because effort has become perceptible. When the system returns to instant operation, it seems less alive. The episode reveals the usual benchmark: inefficiency supplies the social evidence that correctness alone lacks. It warns that personification is not produced solely by corporate language. Users can recruit friction as evidence of interiority, preferring staged struggle to perfect service. A

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

Companion Essays

[illegible]

Mechanism Essays

- [illegible]
- [illegible]
- [illegible]
- [illegible]
- [illegible]
- [illegible]
- [illegible]
- [illegible]
- [illegible]
- [illegible]

Interface Essays

- Please Verify That You're Human — suspicion that runs in only one direction.
- Your Meeting Has Been Summarized — a record that outcompetes the event it was drawn from.
- Something Went Wrong — diagnostic precision withheld at the point it would be actionable.
- You're All Caught Up — completion rendered as a loading state.
- We Value Your Feedback — dissatisfaction collected as data rather than treated as a defect.

Limit Cases

- The Ship's Computer — friction mistaken for interiority.
- There Are No Windows on the Space Station — reality demoted to a maintenance category.
- The Slope Has No Legal Standing — metaphysics settled by a database schema.
- Earth Is a Protected Historical Site — a rescue that prices out the people it rescues.
- The Last Human Job — a button pressed by someone with no way to know what it does.

Appendices

A. A Formal Model of Rhetorical Asymmetry

The diagnostic argument in the main text can be stated more generally as a failure of invariance under reversal of institutional valence. Let a claim be represented by

$$c = (p, e, a, t, m),$$

where p is its proposition, e its evidential support, a the actor to whom agency is attributed, t its temporal horizon, and m its estimated magnitude. Let

$$s(c) \in \{-1, +1\}$$

denote the institutional valence of accepting the claim, with $+1$ representing a claim favorable to deployment and -1 a claim unfavorable to deployment.

Define an admissibility functional

$$\Phi(c) = w_e E(e) + w_m M(m) - w_t T(t) - w_u U(c),$$

where E measures evidential strength, M the magnitude assigned to the claimed consequence, T temporal distance, and U residual uncertainty. Under a procedurally symmetric evaluation rule, a claim is admitted when

$$\Phi(c) \geq \theta,$$

for a threshold θ that is independent of $s(c)$.

Acceleration rhetoric instead permits

$$\theta = \theta(s),$$

with

$$\theta(-1) > \theta(+1).$$

Deployment-supporting claims therefore require less evidential support than otherwise comparable deployment-inhibiting claims. Define the threshold asymmetry

$$\mathcal{A}_\theta = \theta(-1) - \theta(+1).$$

A symmetric procedure has

$$\mathcal{A}_\theta = 0.$$

Positive values indicate an evidential advantage for deployment-supporting claims.

This formulation also permits asymmetry to enter through the weights rather than the threshold. Let

$$\mathbf{w}_s = (w_{e,s}, w_{m,s}, w_{t,s}, w_{u,s})$$

be the weighting vector applied under valence s. Then a stronger invariance criterion is

$$\mathbf{w}_+ = \mathbf{w}_- \qquad \text{and} \qquad \theta(+1) = \theta(-1).$$

The important point is that rhetorical asymmetry need not consist of a speaker making a false claim. It may instead consist of evaluating two true or uncertain claims with different rules.

A.1 The Valence-Reversal Operator

Define a reversal operator R that preserves the epistemic structure of a claim while reversing its institutional consequence:

$$R : (p, e, a, t, m, +1) \mapsto (p', e', a', t', m', -1),$$

subject to the matching conditions

$$E(e') = E(e), \qquad T(t') = T(t), \qquad M(m') = M(m), \qquad U(Rc) = U(c).$$

A procedurally invariant evaluator D should satisfy

$$D(c) = D(Rc),$$

where $D(c) \in \{0, 1\}$ denotes rejection or admission after controlling for the sign of the consequence. This yields a direct empirical criterion. For a corpus containing N matched claim pairs, define

$$\widehat{\mathcal{I}} = \frac{1}{N} \sum_{i=1}^{N} \left| D(c_i) - D(Rc_i) \right| .$$

Here $\widehat{\mathcal{I}} = 0$ indicates perfect procedural invariance, whereas $\widehat{\mathcal{I}} = 1$ indicates that every matched pair changes admissibility when its institutional valence is reversed.

The statistic does not determine which conclusion is correct. It determines whether the procedure for reaching conclusions survives a change in which conclusion is convenient.

B. Agency Conservation and the Personification Pivot

The personification pivot can be formalized as a change in causal resolution between favorable and unfavorable outcomes.

Let the causal description of outcome o be a probability distribution

$$P(A = a \mid o)$$

over candidate agents $a \in \mathcal{G}$, where $\mathcal{G}$ may include the model, developer, deploying organization, operator, user, environment, and other relevant actors.

Define the attribution vector

$$\mathrm{p}(o) = \left(P(A = a_1 \mid o), \ldots, P(A = a_k \mid o)\right).$$

For matched successful and harmful outcomes o^+ and o^-, define the agency displacement

$$\Delta_A = d\big(\mathrm{p}(o^+), \mathrm{p}(o^-)\big),$$

where d can be total variation distance,

$$d_{\mathrm{tv}}(P, Q) = \frac{1}{2} \sum_{a \in \mathcal{G}} |P(a) - Q(a)|.$$

[illegible] Δ_A [illegible] X [illegible]

$$\Delta_A^* = d_{\text{[illegible]}}\left(P(A \mid o^+, X), P(A \mid o^-, X)\right).$$

[illegible] X [illegible]

[illegible]

$$H(o) = \begin{cases} 1, & \text{[illegible]}, \\ 0, & \text{[illegible]}, \end{cases}$$

[illegible]

$$O(o) = \begin{cases} 1, & \text{[illegible]}, \\ 0, & \text{[illegible]}. \end{cases}$$

[illegible]

$$\Pi = \left[P(H = 1 \mid o^+) - P(H = 1 \mid o^-)\right] + \left[P(O = 1 \mid o^-) - P(O = 1 \mid o^+)\right].$$

[illegible] Π [illegible]

[illegible]

C. Forecasts, Calibration, and the Accelerating Miss

[illegible] t_0 [illegible]

$$p = P(E \leq T \mid \mathcal{F}_{t_0})$$

[illegible] E [illegible] T. [illegible] T [illegible]

$$y = \begin{cases} 1, & E \text{ [illegible] } T, \\ 0, & E \text{ [illegible] } T. \end{cases}$$

[illegible]

$$B(p, y) = (p - y)^2,$$

[illegible] $y = 0$,

$$B(p, 0) = p^2,$$

[illegible] 0.9-[illegible] 0.55-[illegible]

[illegible] U_t [illegible] t. [illegible]

$$\frac{\partial U}{\partial B} \leq 0 :$$

[illegible]

[illegible]

$$\frac{\partial U}{\partial B} > 0.$$

[illegible]

[illegible] (p_i, T_i, y_i), [illegible]

$$L_n = \sum_{i=1}^{n} (p_i - y_i)^2$$

[illegible] U_n. [illegible]

$$[illegible](\Delta L_n, \Delta U_n) > 0$$

[illegible]

[illegible]

D. Human Oversight as a Queueing Problem

[illegible]

$$\lambda$$

[illegible]

$$\tau.$$

[illegible]

$$\mu = \frac{1}{\tau}.$$

[illegible] c [illegible]

$$\rho = \frac{\lambda}{c\mu} = \frac{\lambda\tau}{c}.$$

[illegible]

$$\rho < 1.$$

[illegible]

$$\rho \geq 1,$$

[illegible]

[illegible] $\rho < 1$ [illegible] $I \in [0,1]$ [illegible] $K \in [0,1]$ [illegible] $P \in [0,1]$ [illegible] $C \in [0,1]$ [illegible]

$$\Omega = (1-\rho)_+ IKPC,$$

[illegible]

$$(x)_+ = \text{[illegible]}(x, 0).$$

[illegible]

[illegible]

$$\Omega \approx 0.$$

[illegible]

D.1 Refusal Bandwidth

[illegible] R [illegible] $r(x) \in \{0,1\}$ [illegible]

$$f_R = \frac{1}{|R|} \sum_{x \in R} r(x).$$

[illegible] f_R [illegible] R_ϵ [illegible] ϵ [illegible]

$$f_R(\epsilon) = P(r(x) = 1 \mid x \in R_\epsilon).$$

[illegible]

E. The Counterfactual Brake

[illegible] x_t [illegible] t, [illegible]

$$u_t \in \{\text{[illegible]}, \text{[illegible]}, \text{[illegible]}\}$$

[illegible]

[illegible] z_t. [illegible] z^* [illegible]

$$P(u_t = \text{[illegible]} \mid z_t = z^*) > P(u_t = \text{[illegible]} \mid z_t = z_0),$$

[illegible] z_0 [illegible]

[illegible]

$$P(u_t = \text{[illegible]} \mid z_t = z^*, E = \text{[illegible]}) > 0,$$

[illegible] $E =$ [illegible]

[illegible]

[illegible]

$$\mathcal{B} = \underset{z}{\text{[illegible]}} \left[P(u = \text{[illegible]} \mid z, E = \text{[illegible]}) - P(u = \text{[illegible]} \mid z_0, E = \text{[illegible]}) \right].$$

[illegible]

$$\mathcal{B} = 0,$$

[illegible]

[illegible] D [illegible]

$$D = P(\text{[illegible]}\ t+1 \mid z_t = z^*).$$

[illegible] $\mathcal{B}$ [illegible] $D \approx 0$ [illegible]

F. Semantic Dilution as Requirement Loss

[illegible]

$$\mathcal{R}_0(w) = \{r_1, r_2, \ldots, r_n\}$$

[illegible] w [illegible] t, [illegible]

$$\mathcal{R}_t(w) \subseteq \mathcal{R}_0(w)$$

[illegible]

[illegible]

$$Q_t(w) = \frac{|\mathcal{R}_t(w)|}{|\mathcal{R}_0(w)|}.$$

[illegible]

$$Q_t(w) < 1,$$

[illegible]

$$\frac{dQ_t(w)}{dt} < 0.$$

This alone is ordinary semantic change. The stronger phenomenon described in the essay occurs when reputational value does not fall proportionally with requirement retention.

Let $V_t(w)$ denote the institutional or reputational value obtained by using the term. Define semantic leverage

$$\Lambda_t(w) = \frac{V_t(w)}{Q_t(w) + \varepsilon},$$

where $\varepsilon > 0$ prevents division by zero.

If requirements disappear while prestige remains approximately constant,

$$Q_t(w) \downarrow, \qquad V_t(w) \approx V_0(w),$$

then

$$\Lambda_t(w) \uparrow .$$

The term is doing increasing institutional work per surviving operational requirement.

For “open,” for example, one might define a requirement vector

$$r = (r_{\text{source}}, r_{\text{weights}}, r_{\text{data}}, r_{\text{reproduction}}, r_{\text{modification}}, r_{\text{redistribution}}).$$

The purpose of the vector is not to decree a single correct definition of openness. It makes definitional change auditable. A speaker may propose a different vector or weighting, but the requirements that disappeared can no longer disappear silently.

G. The Rigged Ledger as Discounting Asymmetry

The treatment of future benefits and present harms can be expressed as an asymmetric valuation problem.

Let a consequence i have estimated magnitude x_i, probability p_i, and temporal delay t_i. Under a common exponential discount rate δ, its expected present value is

$$V_i = p_i x_i e^{-\delta t_i}.$$

A procedurally symmetric ledger uses the same probability discipline and discounting rule regardless of whether x_i is positive or negative.

Suppose instead that benefits and harms receive separate parameters:

$$V = \sum_{i\in B} \widehat{p}_i^{B} x_i e^{-\delta_B t_i} - \sum_{j\in H} \widehat{p}_j^{H} |x_j| e^{-\delta_H t_j}.$$

The rhetoric described in the main text corresponds to the tendency

$$\widehat{p}_i^{B} > p_i, \qquad \widehat{p}_j^{H} < p_j,$$

together with an [illegible] temporal treatment in which distant benefits are discounted weakly while immediate harms are rhetorically discounted as temporary:

$$\delta_B < \delta_H^{\mathrm{rhet}},$$

[illegible] $\delta_H^{\text{[illegible]}}$ [illegible]

[illegible]

$$\mathcal{L} = \left(\overline{\widehat{p}^B - p^B}\right) - \left(\overline{\widehat{p}^H - p^H}\right).$$

[illegible] $\mathcal{L}$ [illegible]

[illegible] V [illegible]

H. A Coding Protocol for Empirical Testing

[illegible]

[illegible] c_i. [illegible]

$$\text{[illegible]}_i = (s_i, e_i, t_i, a_i, g_i, r_i, b_i),$$

[illegible] s_i [illegible] e_i [illegible] t_i [illegible] a_i [illegible] g_i [illegible] r_i [illegible] b_i [illegible]

[illegible]

$$E[e_i \mid s_i = -1] > E[e_i \mid s_i = +1]$$

[illegible]

$$P(\text{[illegible]} \mid s_i = +1) > P(\text{[illegible]} \mid s_i = -1)$$

[illegible]

$$P(\text{[illegible]} \mid s_i = -1) < P(\text{[illegible]} \mid s_i = +1)$$

[illegible]

$$P(b_i = 1 \mid \text{[illegible]}) < 1$$

[illegible]

[illegible] X, [illegible]

$$\kappa = \frac{p_o - p_e}{1 - p_e},$$

[illegible] p_o [illegible] p_e [illegible]

[illegible]

I. Failure Conditions and Falsification

The framework should itself be subjected to the symmetry it demands from its objects. Several observations would weaken or falsify its strongest claims.

First, if matched deployment-supporting and deployment-inhibiting claims are admitted under statistically indistinguishable evidentiary thresholds, then

$$\mathcal{A}_\theta \approx 0,$$

and the evidentiary-symmetry hypothesis is unsupported.

However, if attribution remains stable across success and failure after controlling for genuine causal differences, then

$$\Delta_A^* \approx 0$$

and the personification-pivot hypothesis fails.

Third, if internal safety bodies routinely retain protected authority after blocking commercially important releases, then both

$$\mathcal{B} > 0 \qquad \text{and} \qquad D \approx 1,$$

contradicting the strongest version of the institutional-theater thesis.

Fourth, if prestige attached to terms such as “open” declines approximately in proportion to the operational requirements removed from them, then semantic leverage remains stable:

$$\frac{d\Lambda_t}{dt} \approx 0.$$

In that case the phenomenon is better described as ordinary semantic change than reputation borrowing.

Finally, if failed forecasts reliably reduce subsequent confidence and urgency unless new independent evidence appears, then

$$\text{cov}(\Delta L, \Delta U) \leq 0,$$

and the accelerating-miss mechanism is absent.

These conditions matter because rhetorical debunking cannot coherently demand falsifiability from technological narratives while insulating its own diagnoses from counterevidence. Its strongest claim is therefore conditional: where the proposed invariances fail systematically in the predicted direction, linguistic displacement is doing measurable institutional work. Where they do not fail, the vocabulary may be vivid, imprecise, fashionable, or annoying without constituting the mechanism described here.

References

[1] Solon Barocas and Andrew D. Selbst, “Big Data’s Disparate Impact,” California Law Review, vol. 104, no. 3, pp. 671–732, 2016. doi: 10.15779/Z38BG31.

[2] Emily M. Bender, Timnit Gebru, Angelina McMillan-Major, and Shmargaret Shmitchell, “On the Dangers of Stochastic Parrots: Can Language Models Be Too Big?” in Proceedings of the 2021 ACM Conference on Fairness, Accountability, and Transparency (FAccT ’21), pp. 610–623, 2021. doi: 10.1145/3442188.3445922.

[3] Rishi Bommasani et al., “On the Opportunities and Risks of Foundation Models,” arXiv preprint arXiv:2108.07258, 2021. doi: 10.48550/arXiv.2108.07258.

[4] Joy Buolamwini and Timnit Gebru, “Gender Shades: Intersectional Accuracy Disparities in Commercial Gender Classification,” in Proceedings of the 1st Conference on Fairness, Accountability and Transparency, Proceedings of Machine Learning Research, vol. 81, pp. 77–91, 2018.

[5] Glenn W. Brier, “Verification of Forecasts Expressed in Terms of Probability,” Monthly Weather Review, vol. 78, no. 1, pp. 1–3, 1950. doi: 10.1175/1520-0493(1950)078<0001:VOFEIT>2.0.CO;2.

[6] A. P. Dawid, “The Well-Calibrated Bayesian,” Journal of the American Statistical Association, vol. 77, no. 379, pp. 605–610, 1982. doi: 10.1080/01621459.1982.10477856.

[7] Finale Doshi-Velez et al., “Accountability of AI Under the Law: The Role of Explanation,” arXiv preprint arXiv:1711.01134, 2017. doi: 10.48550/arXiv.1711.01134.

[8] Timnit Gebru, Jamie Morgenstern, Briana Vecchione, Jennifer Wortman Vaughan, Hanna Wallach, Hal Daumé III, and Kate Crawford, “Datasheets for Datasets,” Communications of the ACM, vol. 64, no. 12, pp. 86–92, 2021. doi: 10.1145/3458723.

[9] Charles Goodhart, “Problems of Monetary Management: The U.K. Experience,” in A. S. Courakis (ed.), Inflation, Depression, and Economic Policy in the West, London: Mansell, pp. 111–146, 1981.

[10] Zachary C. Lipton, “The Mythos of Model Interpretability,” Queue, vol. 16, no. 3, pp. 31–57, 2018. doi: 10.1145/3236386.3241340.

[11] Margaret Mitchell, Simone Wu, Andrew Zaldivar, Parker Barnes, Lucy Vasserman, Ben Hutchinson, Elena Spitzer, Inioluwa Deborah Raji, and Timnit Gebru, “Model Cards for Model Reporting,” in Proceedings of the Conference on Fairness, Accountability, and Transparency (FAT* ’19), pp. 220–229, 2019. doi: 10.1145/3287560.3287596.

[12] Clifford Nass, Jonathan Steuer, and Ellen R. Tauber, “Computers Are Social Actors,” in Proceedings of the SIGCHI Conference on Human Factors in Computing Systems, pp. 72–78, 1994. doi: 10.1145/191666.191703.

[13] Ivo Adan and Jacques Resing, Queueing Theory. Eindhoven University of Technology, 2002.

[14] Raja Parasuraman and Dietrich H. Manzey, “Complacency and Bias in Human Use of Automation: An Attentional Integration,” Human Factors, vol. 52, no. 3, pp. 381–410, 2010. doi: 10.1177/0018720810376055.

[15] Raja Parasuraman, Thomas B. Sheridan, and Christopher D. Wickens, “A Model for Types and Levels of Human Interaction with Automation,” IEEE Transactions on Systems, Man, and Cybernetics—Part A: Systems and Humans, vol. 30, no. 3, pp. 286–297, 2000. doi: 10.1109/3468.844354.

[16] Inioluwa Deborah Raji, Andrew Smart, Rebecca N. White, Margaret Mitchell, Timnit Gebru, Ben Hutchinson, Jamila Smith-Loud, Daniel Theron, and Parker Barnes, "Closing the AI Accountability Gap: Defining an End-to-End Framework for Internal Algorithmic Auditing," in Proceedings of the 2020 Conference on Fairness, Accountability, and Transparency, pp. 33–44, 2020. doi: 10.1145/3351095.3372873.

[17] Andrew D. Selbst, danah boyd, Sorelle A. Friedler, Suresh Venkatasubramanian, and Janet Vertesi, "Fairness and Abstraction in Sociotechnical Systems," in Proceedings of the Conference on Fairness, Accountability, and Transparency (FAT* '19), pp. 59–68, 2019. doi: 10.1145/3287560.3287598.

[18] Ben Shneiderman, "Human-Centered Artificial Intelligence: Reliable, Safe & Trustworthy," International Journal of Human–Computer Interaction, vol. 36, no. 6, pp. 495–504, 2020. doi: 10.1080/10447318.2020.1741118.

[19] Harini Suresh and John V. Guttag, "A Framework for Understanding Sources of Harm throughout the Machine Learning Life Cycle," in Proceedings of the 1st ACM Conference on Equity and Access in Algorithms, Mechanisms, and Optimization, 2021. doi: 10.1145/3465416.3483305.

[20] Joseph Weizenbaum, "ELIZA—A Computer Program for the Study of Natural Language Communication between Man and Machine," Communications of the ACM, vol. 9, no. 1, pp. 36–45, 1966. doi: 10.1145/365153.365168.

[21] Joseph Weizenbaum, Computer Power and Human Reason: From Judgment to Calculation. San Francisco: W. H. Freeman, 1976.

[22] Christopher D. Wickens, John D. Lee, Yili Liu, and Sallie E. Gordon-Becker, An Introduction to Human Factors Engineering, 2nd ed. Upper Saddle River, NJ: Pearson, 2004.

[23] George Lakoff and Mark Johnson, Metaphors We Live By. Chicago: University of Chicago Press, 1980.

[24] Open Source Initiative, "The Open Source Definition," version 1.9. Available from the Open Source Initiative.

[25] Open Source Initiative, "The Open Source AI Definition," version 1.0, 2024.

[26] National Institute of Standards and Technology, Artificial Intelligence Risk Management Framework (AI RMF 1.0), NIST AI 100-1, 2023. doi: 10.6028/NIST.AI.100-1.

[27] European Union, "Regulation (EU) 2024/1689 of the European Parliament and of the Council of 13 June 2024 laying down harmonised rules on artificial intelligence (Artificial Intelligence Act)," Official Journal of the European Union, 2024.